

Telecom Customer Churn machine learning and analysis

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from plotly.offline import plot
%matplotlib inline
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

```
In [2]: data = pd.read_csv('C:\\Users\\windows 10 pro\\Documents\\Telecom+Customer+Churn\\teleco
zipcode = pd.read_csv('C:\\Users\\windows 10 pro\\Documents\\Telecom+Customer+Churn\\tel
```

```
In [9]: data.head()
```

```
Out[9]:
```

	Customer_id	Gender	Age	Married	Number_of_Dependents	City	Zip_Code	Latitude	Longitude	Nu
0	0002-ORFBO	Female	37	Yes	0	Frazier Park	93225	34.827662	-118.999073	
1	0003-MKNFE	Male	46	No	0	Glendale	91206	34.162515	-118.203869	
2	0004-TLHLJ	Male	50	No	0	Costa Mesa	92627	33.645672	-117.922613	
3	0011-IGKFF	Male	78	Yes	0	Martinez	94553	38.014457	-122.115432	
4	0013-EXCHZ	Female	75	Yes	0	Camarillo	93010	34.227846	-119.079903	

5 rows × 38 columns

```
In [4]: zipcode.head()
```

```
Out[4]:
```

	Zip_Code	Population
0	90001	54492
1	90002	44586
2	90003	58198
3	90004	67852
4	90005	43019

```
In [5]: churn_df = pd.merge(data, zipcode[['Zip_Code', 'Population']], on = 'Zip_Code')
```

```
In [10]: churn_df.head()
```

Out[10]:

	Customer_id	Gender	Age	Married	Number_of_Dependents	City	Zip_Code	Latitude	Longitude	Nu
0	0002-ORFBO	Female	37	Yes	0	Frazier Park	93225	34.827662	-118.999073	
1	5183-SNMJQ	Male	32	No	0	Frazier Park	93225	34.827662	-118.999073	
2	6847-KJLTS	Female	72	Yes	0	Frazier Park	93225	34.827662	-118.999073	
3	8788-DOXSU	Male	46	No	0	Frazier Park	93225	34.827662	-118.999073	
4	0003-MKNFE	Male	46	No	0	Glendale	91206	34.162515	-118.203869	

5 rows × 39 columns

In [11]:

churn_df.describe()

Out[11]:

	Age	Number_of_Dependents	Zip_Code	Latitude	Longitude	Number_of_Referrals	Ten
count	7043.000000	7043.000000	7043.000000	7043.000000	7043.000000	7043.000000	
mean	46.509726	0.468692	93486.070567	36.197455	-119.756684	1.951867	
std	16.750352	0.962802	1856.767505	2.468929	2.154425	3.001199	
min	19.000000	0.000000	90001.000000	32.555828	-124.301372	0.000000	
25%	32.000000	0.000000	92101.000000	33.990646	-121.788090	0.000000	
50%	46.000000	0.000000	93518.000000	36.205465	-119.595293	0.000000	
75%	60.000000	0.000000	95329.000000	38.161321	-117.969795	3.000000	
max	80.000000	9.000000	96150.000000	41.962127	-114.192901	11.000000	

In [12]:

churn_df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 7043 entries, 0 to 7042
Data columns (total 39 columns):
```

#	Column	Non-Null Count	Dtype
0	Customer_id	7043 non-null	object
1	Gender	7043 non-null	object
2	Age	7043 non-null	int64
3	Married	7043 non-null	object
4	Number_of_Dependents	7043 non-null	int64
5	City	7043 non-null	object
6	Zip_Code	7043 non-null	int64
7	Latitude	7043 non-null	float64
8	Longitude	7043 non-null	float64
9	Number_of_Referrals	7043 non-null	int64
10	Tenure_in_Months	7043 non-null	int64
11	Offer	7043 non-null	object
12	Phone_Service	7043 non-null	object
13	Avg_Monthly_Long_Distance_Charges	6361 non-null	float64
14	Multiple_Lines	6361 non-null	object
15	Internet_Service	7043 non-null	object
16	Internet_Type	5517 non-null	object
17	Avg_Monthly_GB_Download	5517 non-null	float64
18	Online_Security	5517 non-null	object
19	Online_Backup	5517 non-null	object
20	Device_Protection_Plan	5517 non-null	object
21	Premium_Tech_Support	5517 non-null	object
22	Streaming_TV	5517 non-null	object
23	Streaming_Movies	5517 non-null	object
24	Streaming_Music	5517 non-null	object
25	Unlimited_Data	5517 non-null	object
26	Contract	7043 non-null	object
27	Paperless_Billing	7043 non-null	object
28	Payment_Method	7043 non-null	object
29	Monthly_Charge	7043 non-null	float64
30	Total_Charges	7043 non-null	float64
31	Total_Refunds	7043 non-null	float64
32	Total_Extra_Data_Charges	7043 non-null	int64
33	Total_Long_Distance_Charges	7043 non-null	float64
34	Total_Revenue	7043 non-null	float64
35	Customer_Status	7043 non-null	object
36	Churn_Category	7043 non-null	object
37	Churn_Reason	7043 non-null	object
38	Population	7043 non-null	int64

```
dtypes: float64(9), int64(7), object(23)
```

```
memory usage: 2.1+ MB
```

```
In [13]: churn_df.isna().any()
```

```
Out[13]: Customer_id      False
Gender              False
Age                 False
Married             False
Number_of_Dependents  False
City                False
Zip_Code            False
Latitude            False
Longitude           False
Number_of_Referrals  False
Tenure_in_Months    False
Offer               False
Phone_Service       False
Avg_Monthly_Long_Distance_Charges  True
Multiple_Lines      True
Internet_Service     False
Internet_Type        True
Avg_Monthly_GB_Download  True
Online_Security      True
Online_Backup        True
Device_Protection_Plan  True
Premium_Tech_Support  True
Streaming_TV         True
Streaming_Movies     True
Streaming_Music      True
Unlimited_Data        True
Contract            False
Paperless_Billing    False
Payment_Method       False
Monthly_Charge       False
Total_Charges        False
Total_Refunds        False
Total_Extra_Data_Charges  False
Total_Long_Distance_Charges  False
Total_Revenue        False
Customer_Status      False
Churn_Category       False
Churn_Reason         False
Population           False
dtype: bool
```

EDA and visualization

```
In [14]: cus = churn_df.groupby('Customer_Status')[['Customer_Status']].count()
cus
```

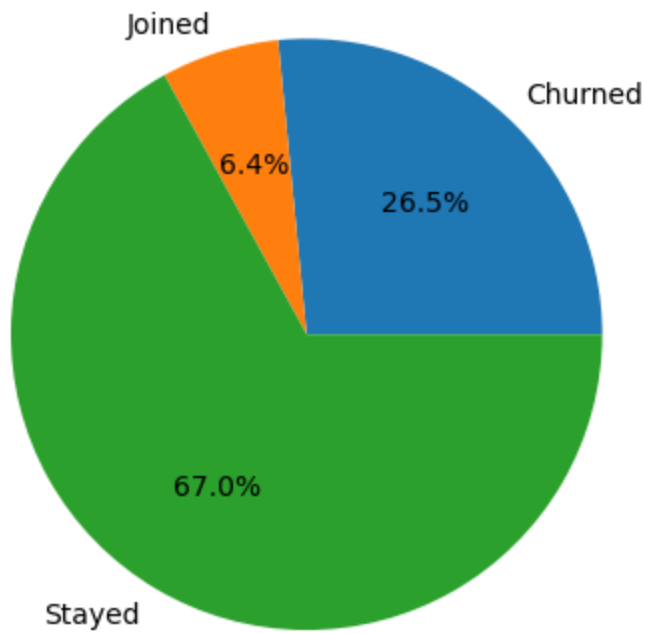
```
Out[14]:
```

Customer_Status	
Customer_Status	
Churned	1869
Joined	454
Stayed	4720

```
In [89]: cus.plot.pie(autopct='%1.1f%%')
plt.title('Customer Churn Status')
plt.ylabel('')
```

```
Out[89]: Text(0, 0.5, '')
```

Customer Churn Status



```
In [113]: churn_total_revenue = churn_df.groupby('Customer_Status')[['Total_Revenue']].sum()  
churn_total_revenue
```


```
Out[113]:
```

	Total_Revenue
Customer_Status	
Churned	3684459.82
Joined	54279.75
Stayed	17632392.12

```
In [112]: churn_total_revenue.plot.pie(autopct='%1.1f%%')  
plt.title('Customer Churn Total Revenue Status')  
plt.ylabel('')
```

```
Out[112]: Text(0, 0.5, '')
```

Customer Churn Total Revenue Status



```
In [120]: churn_monthly_charge = churn_df.groupby('Customer_Status')[['Monthly_Charge']].sum()  
churn_monthly_charge
```

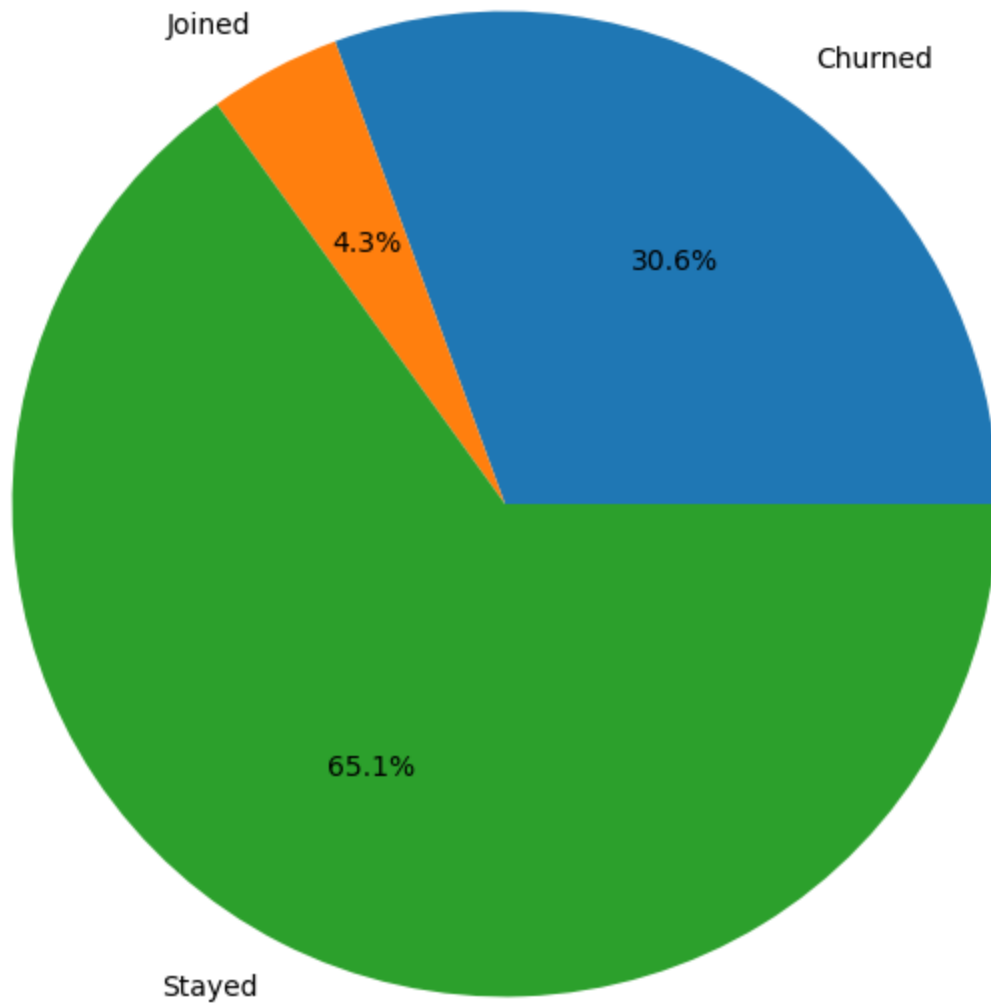
```
Out[120]:
```

	Monthly_Charge
Customer_Status	
Churned	137086.65
Joined	19420.30
Stayed	291400.60

```
In [119]: churn_monthly_charge.plot.pie(autopct='%1.1f%%')  
plt.title('Customer Churn Monthly Charge Status')  
plt.ylabel('')
```

```
Out[119]: Text(0, 0.5, '')
```

Customer Churn Monthly Charge Status



```
In [124]: churn_total_long_distance_charges = churn_df.groupby('Customer_Status')[['Total_Long_Dis  
churn_total_long_distance_charges
```


```
Out[124]:
```

Total_Long_Distance_Charges	
Customer_Status	
Churned	810991.9
Joined	17309.2
Stayed	4447605.0

```
In [123]: churn_total_long_distance_charges.plot.pie(autopct='%1.1f%%')  
plt.title('Customer Churn Total Long Distance Charges')  
plt.ylabel('')
```

```
Out[123]: Text(0, 0.5, '')
```

Customer Churn Total Long Distance Charges



```
In [92]: churn_cat = churn_df.groupby('Churn_Category')[['Churn_Category']].count()  
churn_cat
```

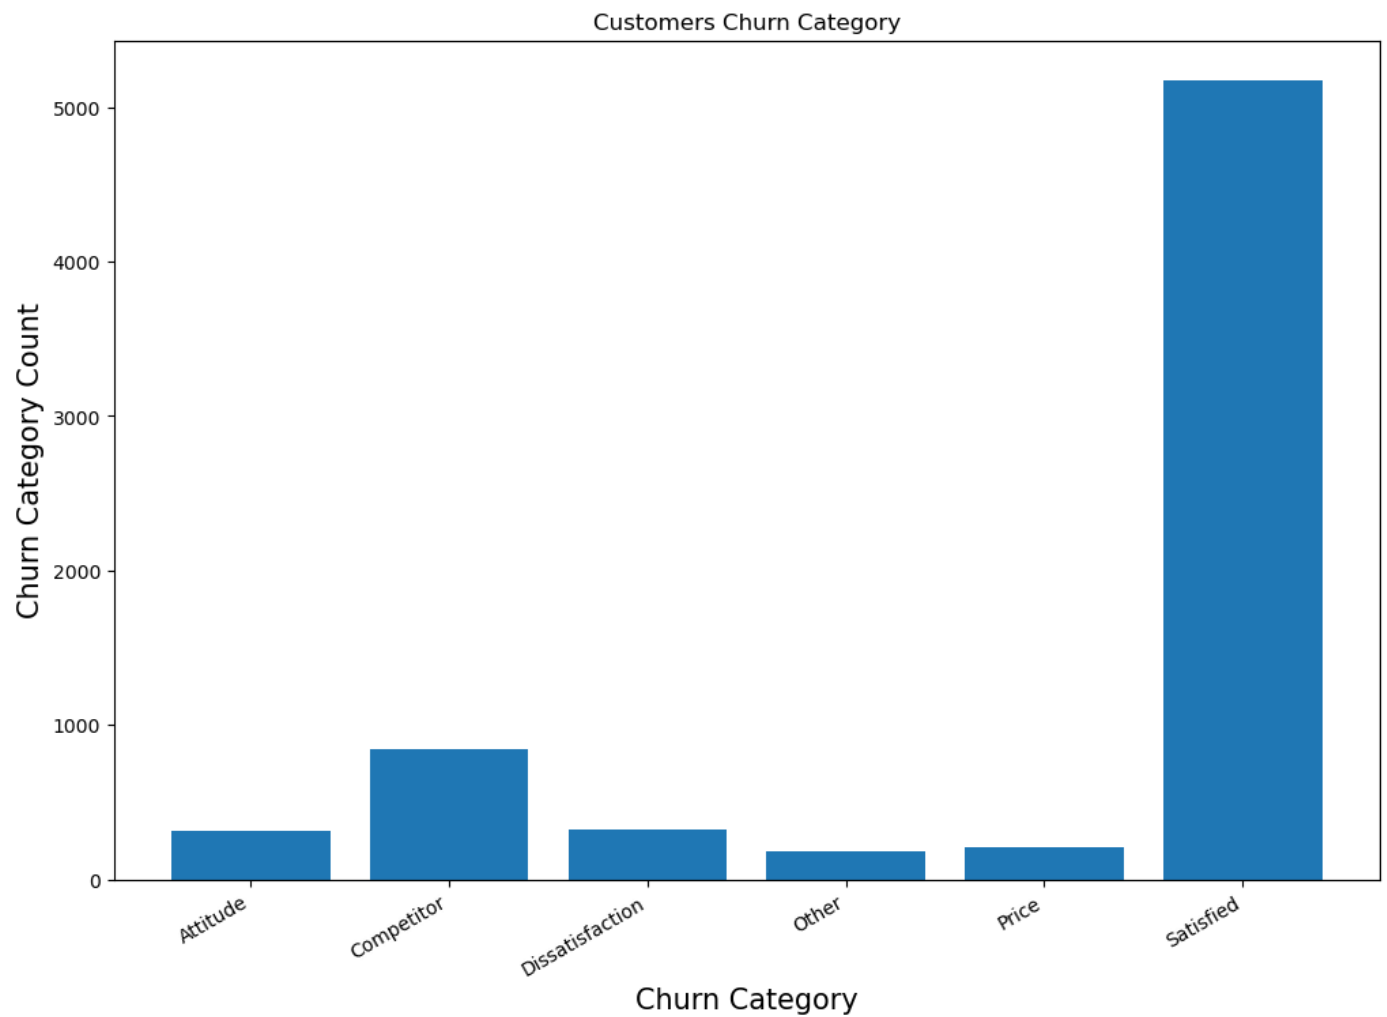
```
Out[92]:
```

Churn_Category	
----------------	--

Churn_Category	
Attitude	314
Competitor	841
Dissatisfaction	321
Other	182
Price	211
Satisfied	5174

```
In [97]: plt.rcParams['figure.figsize']=(12,8)  
plt.title('Customers Churn Category')  
plt.xlabel('Churn Category', fontsize = 15)  
plt.ylabel('Churn Category Count', fontsize = 15)  
plt.xticks(rotation = 30, ha = 'right')  
plt.bar(churn_cat.index, churn_cat.Churn_Category)
```


Out[97]: <BarContainer object of 6 artists>



```
In [102... churn_rea = churn_df.groupby('Churn_Reason')[['Churn_Reason']].count()  
churn_rea
```


Out[102]:

Churn_Reason

Churn_Reason	
Attitude of service provider	94
Attitude of support person	220
Competitor had better devices	313
Competitor made better offer	311
Competitor offered higher download speeds	100
Competitor offered more data	117
Deceased	6
Don't know	130
Extra data charges	39
Lack of affordable download/upload speed	30
Lack of self-service on Website	29
Limited range of services	37
Long distance charges	64
Moved	46
Network reliability	72
Poor expertise of online support	31
Poor expertise of phone support	12
Price too high	78
Product dissatisfaction	77
Satisfied	5174
Service dissatisfaction	63

```
In [103]: plt.rcParams['figure.figsize']=(12,8)
plt.title('Customers Churn Reason')
plt.xlabel('Churn Reason', fontsize = 15)
plt.ylabel('Churn Reason Count', fontsize = 15)
plt.xticks(rotation = 30, ha = 'right')
plt.bar(churn_rea.index, churn_rea.Churn_Reason)
```

Out[103]: <BarContainer object of 21 artists>



```
In [132]: # Total closing rate by Region
total_City = churn_df.groupby('City')[['Total_Revenue']].sum()
total_City
```

Out[132]:

Total_Revenue	
City	
Acampo	18107.96
Acton	12156.36
Adelanto	18235.49
Adin	5539.38
Agoura Hills	10641.88
...	...
Yreka	17467.67
Yuba City	17867.82
Yucaipa	29765.80
Yucca Valley	12374.69
Zenia	12733.38

1106 rows × 1 columns

```
In [134... top_10 = total_City.sort_values('Total_Revenue', ascending=False).head(10)
print(top_10)
bottom_10 = total_City.sort_values('Total_Revenue', ascending=False).tail(10)
print(bottom_10)
```

City	Total_Revenue
Los Angeles	852725.23
San Diego	738416.01
Sacramento	353371.84
San Jose	326478.36
San Francisco	306995.99
Fresno	194430.25
Long Beach	185937.12
Escondido	155899.80
Oakland	154564.36
Whittier	128858.77

City	Total_Revenue
Selma	1613.53
Twain	1453.58
Eldridge	1405.04
Homeland	1400.54
Arnold	1215.61
Highland	1193.64
Loleta	873.05
Angelus Oaks	744.72
Dana Point	694.86
Eagleville	238.57


```
In [135... # Total Revenue by City Visualization
fig, axes = plt.subplots (1,2, figsize = (16,6))
plt.tight_layout(pad=2)
xlabels = top_10.index
axes[0].set_title('Top 10 city with highest revenue')
axes[0].set_xticklabels(xlabels, rotation = 45, ha = 'right')
sns.barplot(x=top_10.index, y=top_10.Total_Revenue, ax=axes[0])
axes[0].set_xlabel('City')
axes[0].set_ylabel('Total_Revenue')

xlabels = bottom_10.index
axes[1].set_title('Bottom 10 city with lowest revenue')
axes[1].set_xticklabels(xlabels, rotation = 45, ha = 'right')
sns.barplot(x=bottom_10.index, y=bottom_10.Total_Revenue, ax=axes[1])
axes[1].set_xlabel('City')
axes[1].set_ylabel('Total_Revenue')
```

C:\Users\windows 10 pro\AppData\Local\Temp\ipykernel_5232\906439376.py:6: UserWarning:
FixedFormatter should only be used together with FixedLocator

C:\Users\windows 10 pro\AppData\Local\Temp\ipykernel_5232\906439376.py:13: UserWarning:
FixedFormatter should only be used together with FixedLocator

```
Out[135]: Text(838.9608585858584, 0.5, 'Total_Revenue')
```



```
In [15]: churn_df1 = churn_df.iloc[:, [13,17,29,30,31,32,33,35]]
churn_df1.head(5)
```

	Avg_Monthly_Long_Distance_Charges	Avg_Monthly_GB_Download	Monthly_Charge	Total_Charges	Total_Refu
0	42.39	16.0	65.60	593.3	
1	45.69	11.0	95.10	865.1	4
2	47.34	28.0	100.40	5749.8	
3	9.70	6.0	61.35	3645.5	
4	10.69	10.0	-4.00	542.4	3

```
In [16]: churn_df1['Avg_Monthly_Long_Distance_Charges'].fillna(churn_df1['Avg_Monthly_Long_Distance_Charges'].mean(), inplace=True)
churn_df1['Avg_Monthly_GB_Download'].fillna(churn_df1['Avg_Monthly_GB_Download'].mean(), inplace=True)
```

C:\Users\windows 10 pro\AppData\Local\Temp\ipykernel_1124\74051609.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

C:\Users\windows 10 pro\AppData\Local\Temp\ipykernel_1124\74051609.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
In [17]: churn_df1.isna().any()
```

```
Out[17]: Avg_Monthly_Long_Distance_Charges    False
Avg_Monthly_GB_Download                 False
Monthly_Charge                         False
Total_Charges                          False
Total_Refunds                          False
Total_Extra_Data_Charges                False
Total_Long_Distance_Charges             False
Customer_Status                        False
dtype: bool
```

```
In [18]: churn_df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 7043 entries, 0 to 7042
Data columns (total 8 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Avg_Monthly_Long_Distance_Charges    7043 non-null  float64
 1   Avg_Monthly_GB_Download               7043 non-null  float64
 2   Monthly_Charge                       7043 non-null  float64
 3   Total_Charges                        7043 non-null  float64
 4   Total_Refunds                        7043 non-null  float64
 5   Total_Extra_Data_Charges             7043 non-null  int64
 6   Total_Long_Distance_Charges          7043 non-null  float64
 7   Customer_Status                      7043 non-null  object
dtypes: float64(6), int64(1), object(1)
memory usage: 495.2+ KB
```

```
In [19]: X = churn_df1.drop('Customer_Status', axis=1)
y = churn_df1['Customer_Status']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
```

```
In [ ]:
```

```
In [20]: # Feature scaling
sc = MinMaxScaler(feature_range=(0,1))
sc1 = sc.fit_transform(X_train)
```

```
In [21]: clf1 = RandomForestClassifier(n_jobs=3, random_state=0)
clf1.fit(X_train, y_train)
```

```
Out[21]: ▼ Random Forest Classifier
RandomForestClassifier(n_jobs=3, random_state=0)
```

```
In [22]: # Make predictions
y_pred = clf1.predict(X_test)
```

```
In [23]: accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")

Accuracy: 0.7885024840312278
```

```
In [35]: con_matrix = confusion_matrix(y_test, y_pred)
print(f"Accuracy: {con_matrix}")

Accuracy: [[112  30 234]
 [ 46  40   8]
 [ 25   0 914]]
```

```
In [24]: report = classification_report(y_test, y_pred)
print(f"Accuracy: {report}")
```

Accuracy:		precision	recall	f1-score	support
Churned	0.65	0.48	0.55	376	
Joined	0.58	0.49	0.53	94	
Stayed	0.84	0.94	0.89	939	
accuracy			0.79	1409	
macro avg	0.69	0.64	0.66	1409	
weighted avg	0.77	0.79	0.77	1409	

Application of 20% hypothesis discount on the charge

```
In [25]: X_test_20_percent_discount = X_test
X_test_20_percent_discount.head()
```

	Avg_Monthly_Long_Distance_Charges	Avg_Monthly_GB_Download	Monthly_Charge	Total_Charges	Total_F
2200	23.13	24.0	109.90	7332.40	
4627	1.15	23.0	70.10	70.10	
3225	13.40	17.0	44.00	44.00	
2828	33.69	15.0	80.00	1029.35	
3768	1.38	16.0	69.95	320.40	

```
In [26]: X_test_20_percent_discount.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1409 entries, 2200 to 3065
Data columns (total 7 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Avg_Monthly_Long_Distance_Charges    1409 non-null   float64
1   Avg_Monthly_GB_Download               1409 non-null   float64
2   Monthly_Charge                       1409 non-null   float64
3   Total_Charges                        1409 non-null   float64
4   Total_Refunds                        1409 non-null   float64
5   Total_Extra_Data_Charges             1409 non-null   int64
6   Total_Long_Distance_Charges          1409 non-null   float64
dtypes: float64(6), int64(1)
memory usage: 88.1 KB
```

The 20% discount hypothesis is applied using X_test column Monthly_Charge and Total_Long_Distance_Charges

```
In [27]: X_test_20_percent_discount['Monthly_Charge'] = X_test_20_percent_discount['Monthly_Charg
X_test_20_percent_discount['Total_Long_Distance_Charges'] = X_test_20_percent_discount['
```

```
In [28]: X_test_20_percent_discount.head()
```

	Avg_Monthly_Long_Distance_Charges	Avg_Monthly_GB_Download	Monthly_Charge	Total_Charges	Total_F
2200	23.13	24.0	87.92	7332.40	
4627	1.15	23.0	56.08	70.10	
3225	13.40	17.0	35.20	44.00	
2828	33.69	15.0	64.00	1029.35	
3768	1.38	16.0	55.96	320.40	

```
In [39]: # Make predictions
y_pred_20 = clf1.predict(X_test_20_percent_discount)
```

```
In [40]: accuracy = accuracy_score(y_test, y_pred_20)
print(f"Accuracy: {accuracy}")
```

Accuracy: 0.7565649396735273

```
In [41]: con_matrix1 = confusion_matrix(y_test, y_pred_20)
print(f"Accuracy: {con_matrix}")
```

Accuracy: [[112 30 234]
 [46 40 8]
 [25 0 914]]

```
In [42]: report = classification_report(y_test, y_pred_20)
print(f"Accuracy: {report}")
```

Accuracy:		precision	recall	f1-score	support
Churned	0.61	0.30	0.40	0.37	376
Joined	0.57	0.43	0.49	0.51	94
Stayed	0.79	0.97	0.87	0.92	939
accuracy			0.76		1409
macro avg	0.66	0.57	0.59		1409
weighted avg	0.73	0.76	0.72		1409

From the above before and after 20% hypothesis discount, the classification report show that the support in the model

is same. This proof that the 20% do not have any impact on churn.

Findings and Recommendations

churn is indeed high in the internet division

* 25.5% across 1869 out of 7043 customers

Predictive model is able to predict churn but the main driver is not customer charge sensitivity

* Attitude of support person, Competitor had better devices, and Competitor made better offer are the largest drivers

Discount strategy of 20% discount is very less effective and not necessary.

* Application of 20% discount can be used to motivate new subscriber to join

```
In [ ]:
```